

# Round Robin Server Load Balancer System - Report

## Distributed Server Load Balancer System

### Team Members

Mohamed Safwat Hassan - 192100140

Abdallah Abdelmonem Abdelaziz - 192100154

Abdelrahman Omar Ali - 191900004

Hamed Mohamed Ragab - 191900056

Ahmed Galal Abdelhay - 192200374

### Technical Report

#### Executive Summary

This report details the architecture and implementation of a distributed server load balancer system built using Flask for the backend and HTML/CSS/JavaScript for the frontend. The system monitors multiple backend servers, distributes incoming traffic among them using a round-robin algorithm, tracks their health status, and visualizes performance metrics through an intuitive dashboard interface.

The project demonstrates key concepts in distributed systems, including:

- Load balancing
- Health monitoring
- Performance metrics collection and visualization
- Fault tolerance
- Real-time data processing

#### System Architecture

The distributed server load balancer consists of several interconnected components:

1. **Load Balancer** ( `load_balancer.py` ): The central component that routes incoming requests to available backend servers using a round-robin approach.
2. **Backend Servers** ( `server1.py` , `server2.py` , `server3.py` , `server4.py` ): Multiple identical servers running on different ports that handle client requests.
3. **Dashboard** ( `dashboard.py` ): A web application that provides real-time monitoring and visualization of the system status.
4. **Testing Module** ( `testing_servers.py` ): A script that generates test traffic to demonstrate the load balancer's functionality.
5. **Frontend** ( `index.html` , `script.js` , `styles.css` ): The user interface for the dashboard, displaying server metrics and visualizations.

# Component Analysis

## Load Balancer ( `load_balancer.py` )

The load balancer serves as the entry point for all client requests and implements several key features:

- **Round-Robin Scheduling:** Uses Python's `itertools.cycle` to distribute incoming requests evenly across available servers.
- **Health Monitoring:** Continuously checks the health of backend servers through a dedicated thread.
- **Request Tracking:** Counts requests sent to each server for load analysis.
- **Metrics Collection:** Gathers and stores performance data for each server, including:
  - CPU usage
  - RAM usage
  - Response latency
  - Uptime
  - Request count
- **API Endpoints:**
  - `/` : Redirects incoming requests to available backend servers
  - `/stats` : Provides detailed metrics and status information for all servers

The load balancer also implements Cross-Origin Resource Sharing (CORS) to allow the dashboard to fetch data from its API.

## Backend Servers ( `serverX.py` )

Each backend server is a Flask application running on a different port (5001-5004) and offers the following endpoints:

- `/` : Returns a simple greeting message
- `/health` : Simple endpoint for health checks, returns "OK" with status code 200
- `/metrics` : Returns detailed server metrics:
  - CPU usage (percentage)
  - RAM usage (MB)
  - Uptime (formatted as HH:MM:SS)

The servers use the `psutil` library to collect system metrics for the specific Python process.

## Dashboard ( `dashboard.py` )

The dashboard provides a web interface to monitor the status of all backend servers. It includes:

- **API Integration:** Fetches stats from all backend servers and aggregates them
- **Error Handling:** Gracefully handles server failures and timeouts
- **Data Aggregation:** Collects and formats metrics from all servers for the frontend

## Testing Module ( `testing_servers.py` )

A simple Python script that generates continuous traffic to the load balancer for testing purposes:

- Sends 100 requests with a 2-second delay between requests
- Repeats this process 50 times
- Outputs the status code and response from each request

## Frontend Components

### HTML Structure ( `index.html` ):

- Responsive layout with separate sections for different visualizations
- Server status table with dynamic data
- Charts container for visualizing request distribution and server health
- Interactive hover card with detailed server information

### JavaScript ( `script.js` ):

- Fetches server statistics every 2 seconds
- Updates the UI elements in real-time
- Implements three types of visualizations:

1. **Server Status Table:** Shows basic information with color-coded status indicators
2. **Request Distribution Pie Chart:** Visualizes the proportion of requests handled by each server
3. **Health/Requests Line Graphs:** Tracks server health and request count over time

- Interactive elements:

- Hover functionality for detailed server information
- Tooltip positioning that adapts to screen boundaries

## Data Flow

### 1. Client Request Flow:

- Client sends request to the load balancer (port 5000)
- Load balancer selects the next available server using round-robin algorithm
- Request is redirected to the selected backend server (ports 5001-5004)
- Backend server processes the request and returns a response
- Response is returned to the client

### 2. Monitoring Flow:

- Dashboard (port 8000) fetches stats from the load balancer's `/stats` endpoint
- Load balancer gathers health and metrics data from each backend server
- Dashboard processes and visualizes this data through the frontend
- Frontend updates every 2 seconds with fresh data

## Key Features

### 1. Fault Tolerance:

- The load balancer detects server failures through regular health checks
- Failed servers are excluded from request distribution
- The system continues operating with reduced capacity when servers fail

## 2. Dynamic Server Health Monitoring:

- Continuous background health checks via a dedicated thread
- Real-time updates of server status (UP/DOWN)
- Visualization of server health over time

## 3. Performance Metrics:

- CPU and RAM usage tracking
- Response latency measurement
- Request distribution analysis
- Server uptime monitoring

## 4. Interactive Dashboard:

- Real-time data updates without page refresh
- Multiple visualization types for different metrics
- Detailed server information on hover
- Color-coded status indicators

## 5. Historical Data:

- Time-series graphs showing server health and requests over time
- JSON files stored in `api/stats/` for potential historical analysis

# Implementation Details

## Round-Robin Load Balancing

The system implements a classic round-robin load balancing algorithm using Python's `itertools.cycle`:

```
server_cycle = itertools.cycle(servers)
```

When a request arrives, the load balancer selects the next server in the rotation with this code:

```
target = next(server_cycle)

if server_status.get(target, False):

    server_request_count[target] += 1
```

```
return redirect(f"{target}/", code=307)
```

This ensures even distribution of requests among healthy servers.

## Health Monitoring

The health check function runs in a separate thread:

```
def health_check():  
  
    while True:  
  
        for server in servers:  
  
            try:  
  
                response = requests.get(server + "/health", timeout=1)  
  
                server_status[server] = (response.status_code == 200)  
  
            except requests.exceptions.RequestException:  
  
                server_status[server] = False  
  
        time.sleep(2)
```

This allows the load balancer to maintain an up-to-date view of server availability without blocking the main request handling thread.

## Data Visualization

The frontend implements multiple visualization techniques:

1. **Server Status Table** with visual indicators:

- Green/red dot for UP/DOWN status
- Server URL and request count

2. **Interactive Pie Chart** showing request distribution:

- Color-coded segments for each server
- Consistent colors across refreshes (using URL hash)
- Interactive legends and tooltips

3. **Time-Series Line Graphs** for each server:

- Dual-axis chart showing both requests (left axis) and health status (right axis)
- Color-coded lines (green for requests, red for health)
- Real-time updates with sliding window (last 30 data points)

## Testing and Results

The `testing_servers.py` script simulates client traffic by sending 5,000 requests (100 requests × 50 iterations) to the load balancer. With the round-robin algorithm, we expect each server to handle approximately 1,250 requests if all servers remain healthy.

In practice, the distribution may vary slightly due to:

- Server health status changes
- Response time variations
- Network conditions

The dashboard allows visualization of this distribution in real-time, confirming the effectiveness of the load balancing strategy.

## Limitations and Future Improvements

### 1. Load Balancing Algorithm:

- The current implementation uses simple round-robin without considering server load
- Future improvement: Implement weighted round-robin or least-connections algorithms

### 2. Persistence:

- The system does not maintain session persistence
- Future improvement: Add support for sticky sessions with cookie-based routing

### 3. Security:

- Basic implementation without authentication or HTTPS
- Future improvement: Add TLS/SSL support and authentication mechanisms

### 4. Scalability:

- Fixed list of backend servers
- Future improvement: Implement dynamic server discovery and auto-scaling

### 5. High Availability:

- Single load balancer instance creates a single point of failure
- Future improvement: Implement redundant load balancers with failover capability

## Conclusion

This distributed server load balancer project successfully demonstrates core principles of distributed systems and load balancing. The combination of Flask-based microservices with a responsive web dashboard provides both functionality and visibility into system operations.

The implementation showcases:

- Effective request distribution across multiple servers
- Real-time monitoring and visualization
- Fault tolerance through health checking
- Performance metrics collection and analysis

While there are several areas for potential improvement, the current system provides a solid foundation for understanding distributed system concepts and load balancing techniques.

## Appendix: Project Structure

```
project/

├─ api/

|   └─ stats/

|       └─ [json files for servers status]

├─ static/

|   └─ scripts/

|       └─ script.js

|   └─ styles/

|       └─ styles.css

├─ templates/

|   └─ index.html

├─ dashboard.py

├─ load_balancer.py

├─ testing_servers.py

└─ serverX.py (4 server instances)
```